

✓ Student Information

Name: Rohan Agrawal

PRN: 202201040208

Batch: Batch 1

LOGISTIC REGRESSION (USING SKLEARN)

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, roc_curve, auc

# Load the CSV file (Ensure the correct path is given)
df = pd.read_csv("heart1.csv")

# Assume the last column is the target variable and the rest are features
X = df.iloc[:, :-1].values # Features
y = df.iloc[:, -1].values # Target variable

# Split the dataset into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardizing features for better model performance
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Create and train the Logistic Regression model
model = LogisticRegression()
model.fit(X_train, y_train)
```

```
# Predict on the test data
y_pred = model.predict(X_test)
y_pred_proba = model.predict_proba(X_test)[:, 1] # Probabilities for ROC curve



# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy using Scikit-Learn: {accuracy:.4f}")

# 1. Confusion Matrix
plt.figure(figsize=(6, 4))
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["No Disease", "Disease"], yticklabels=["No Disease", "Disease"])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()

# 2. ROC Curve
# Plots True Positive Rate (TPR) vs False Positive Rate (FPR).
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(6, 4))
plt.plot(fpr, tpr, color="blue", label=f"ROC Curve (AUC = {roc_auc:.2f})")
plt.plot([0, 1], [0, 1], linestyle="--", color="red") # Diagonal baseline
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.legend()
plt.show()
```

↔ Accuracy using Scikit-Learn: 0.7951



LOGISTIC REGRESSION (SCRATCH)

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, roc_curve, auc

# Sigmoid function to map values to probabilities
def sigmoid(z):
    return 1 / (1 + np.exp(-z))

# Function to compute cost (log loss)
def compute_cost(X, y, weights):
    m = len(y)
    predictions = sigmoid(np.dot(X, weights))
    cost = (-1/m) * np.sum(y * np.log(predictions) + (1 - y) * np.log(1 - predictions))
    return cost

# Function to perform gradient descent
def gradient_descent(X, y, weights, learning_rate, iterations):
    m = len(y)
    cost_history = []

    for i in range(iterations):
        predictions = sigmoid(np.dot(X, weights))
        error = predictions - y
        gradient = (1/m) * np.dot(X.T, error)
        weights -= learning_rate * gradient

        cost = compute_cost(X, y, weights)
        cost_history.append(cost)

        if i % 1000 == 0:
            print(f"Iteration {i}: Cost {cost:.4f}")

    return weights, cost_history

# Load the dataset
df = pd.read_csv("heart1.csv")

# Assume the last column is the target variable and the rest are features
```



```
X = df.iloc[:, :-1].values # Features
y = df.iloc[:, -1].values # Target variable

# Split data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardizing the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Add a bias term (column of ones) to X_train and X_test
X_train = np.c_[np.ones((X_train.shape[0], 1)), X_train]
X_test = np.c_[np.ones((X_test.shape[0], 1)), X_test]

# Initialize weights with zeros
weights = np.zeros(X_train.shape[1])

# Hyperparameters
learning_rate = 0.01
iterations = 10000

# Train the model
weights, cost_history = gradient_descent(X_train, y_train, weights, learning_rate, iterations)

# Make predictions
y_pred_proba = sigmoid(np.dot(X_test, weights)) # Probabilities for ROC curve
y_pred = y_pred_proba >= 0.5 # Convert probabilities to 0 or 1

# Calculate accuracy
accuracy = np.mean(y_pred == y_test)
print(f"Accuracy from scratch: {accuracy:.4f}")



# 1. Confusion Matrix
plt.figure(figsize=(6, 4))
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["No Disease", "Disease"], yticklabels=["No Disease", "Disease"])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```


```
# 2. ROC Curve
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(6, 4))
plt.plot(fpr, tpr, color="blue", label=f"ROC Curve (AUC = {roc_auc:.2f})")
plt.plot([0, 1], [0, 1], linestyle="--", color="red") # Diagonal baseline
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.legend()
plt.show()
```



```
↵ Iteration 0: Cost 0.6898  
Iteration 1000: Cost 0.3453  
Iteration 2000: Cost 0.3339  
Iteration 3000: Cost 0.3311  
Iteration 4000: Cost 0.3301  
Iteration 5000: Cost 0.3298  
Iteration 6000: Cost 0.3296  
Iteration 7000: Cost 0.3295  
Iteration 8000: Cost 0.3295  
Iteration 9000: Cost 0.3295  
Accuracy from scratch: 0.7951
```





LOGLOSS (SCRATCH)

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, roc_curve, auc

# Sigmoid function
def sigmoid(z):
    return 1 / (1 + np.exp(-z))

# Log Loss function
def compute_log_loss(y_true, y_pred_proba):
    m = len(y_true)
    epsilon = 1e-15 # Avoid log(0)
    y_pred_proba = np.clip(y_pred_proba, epsilon, 1 - epsilon)
    loss = -np.mean(y_true * np.log(y_pred_proba) + (1 - y_true) * np.log(1 - y_pred_proba))
    return loss

# Gradient Descent function
def gradient_descent(X, y, weights, learning_rate, iterations):
    m = len(y)
    log_loss_history = []

    for i in range(iterations):
        predictions = sigmoid(np.dot(X, weights))
        error = predictions - y
        gradient = (1/m) * np.dot(X.T, error)
        weights -= learning_rate * gradient

        # Compute and store log loss
        loss = compute_log_loss(y, predictions)
        log_loss_history.append(loss)

        if i % 1000 == 0:
            print(f"Iteration {i}: Log Loss = {loss:.4f}")

    return weights, log_loss_history

# Load the dataset
df = pd.read_csv("heart1.csv")
```



```
# Assume the last column is the target variable and the rest are features
X = df.iloc[:, :-1].values # Features
y = df.iloc[:, -1].values # Target variable

# Split data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardizing the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Add bias term (column of ones)
X_train = np.c_[np.ones((X_train.shape[0], 1)), X_train]
X_test = np.c_[np.ones((X_test.shape[0], 1)), X_test]

# Initialize weights with zeros
weights = np.zeros(X_train.shape[1])

# Hyperparameters
learning_rate = 0.01
iterations = 10000

# Train model
weights, log_loss_history = gradient_descent(X_train, y_train, weights, learning_rate, iterations)

# Predictions
y_pred_proba = sigmoid(np.dot(X_test, weights))
y_pred = y_pred_proba >= 0.5 # Convert to 0 or 1


# Compute Accuracy
accuracy = np.mean(y_pred == y_test)
print(f"Final Accuracy: {accuracy:.4f}")

# ----- Plot Graphs -----

# 1. Log Loss Curve (Cost Function)
plt.figure(figsize=(6, 4))
plt.plot(range(len(log_loss_history)), log_loss_history, color="blue", label="Log Loss")
plt.xlabel("Iterations")
plt.ylabel("Log Loss")
```

```
plt.title("Log Loss Convergence Over Iterations")  
plt.legend()  
plt.grid()  
plt.show()
```

↔ Iteration 0: Log Loss = 0.6931
Iteration 1000: Log Loss = 0.3453
Iteration 2000: Log Loss = 0.3339
Iteration 3000: Log Loss = 0.3311
Iteration 4000: Log Loss = 0.3301
Iteration 5000: Log Loss = 0.3298
Iteration 6000: Log Loss = 0.3296
Iteration 7000: Log Loss = 0.3295
Iteration 8000: Log Loss = 0.3295
Iteration 9000: Log Loss = 0.3295
Final Accuracy: 0.7951



All Activation Functions

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler

# Load the Heart dataset
df = pd.read_csv("heart1.csv") # Replace with actual file path or URL

# Fill missing values (if any) in a column of interest, e.g., 'Age' or 'Cholesterol'
df['age'] = df['age'].fillna(df['age'].median()) # Example: Replace missing values in 'Age'

# Select the 'Age' column (or any other numeric column of interest) as our feature (X)
X = df['age'].values

# Standardize the feature (important for activation functions)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X.reshape(-1, 1)).flatten() # Reshape and standardize

# Apply activation functions on the feature (ReLU, Sigmoid, Tanh)
y_relu = np.maximum(0, X_scaled) # ReLU
y_sigmoid = 1 / (1 + np.exp(-X_scaled)) # Sigmoid
y_tanh = np.tanh(X_scaled) # Tanh

# Create subplots to display the activation functions
plt.figure(figsize=(12, 8))

# ReLU plot
plt.subplot(3, 1, 1)
plt.plot(X_scaled, y_relu, label="ReLU", color='blue', marker='o', linestyle='none', markersize=4)
plt.title("ReLU Activation Function on Age")
plt.xlabel("Scaled Age")
plt.ylabel("ReLU(Scaled Age)")
plt.grid(True)
plt.legend()

# Sigmoid plot
plt.subplot(3, 1, 2)
plt.plot(X_scaled, y_sigmoid, label="Sigmoid", color='green', marker='x', linestyle='none', markersize=4)
plt.title("Sigmoid Activation Function on Age")
plt.xlabel("Scaled Age")
plt.ylabel("Sigmoid(Scaled Age)")
plt.grid(True)
```


```
plt.legend()

# Tanh plot
plt.subplot(3, 1, 3)
plt.plot(X_scaled, y_tanh, label="Tanh", color='red', marker='s', linestyle='none', markersize=4)
plt.title("Tanh Activation Function on Age")
plt.xlabel("Scaled Age")
plt.ylabel("Tanh(Scaled Age)")
plt.grid(True)
plt.legend()


# Show the plots
plt.tight_layout()
plt.show()
```





ReLU Activation Function on Age



Sigmoid Activation Function on Age



Tanh Activation Function on Age



ANN using sklearn

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import seaborn as sns

# Load the Heart dataset
df = pd.read_csv("heart1.csv") # Replace with actual file path or URL

# Preprocess the data (e.g., fill missing values, select features, etc.)
df.fillna(df.median(), inplace=True) # Fill missing values with median (for simplicity)

# Select features and target (assuming 'target' is the column with labels)
X = df.drop('target', axis=1).values # Features (adjust column name)
y = df['target'].values # Target

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Standardize the features (important for neural networks)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Initialize and train the MLPClassifier (ANN)
mlp = MLPClassifier(hidden_layer_sizes=(30,), max_iter=1000, activation='relu', random_state=42)
mlp.fit(X_train_scaled, y_train)

# Make predictions
y_pred = mlp.predict(X_test_scaled)

# Accuracy and classification report
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy * 100:.2f}%")
print("Classification Report:\n", classification_report(y_test, y_pred))

# Confusion Matrix plot
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 5))
```

```
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=['No Disease', 'Disease'], yticklabels=['No Disease', 'Disease'])
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
# Plot the loss curve during training
plt.figure(figsize=(8, 6))
plt.plot(mlp.loss_curve_, label='Training Loss')
plt.title('Training Loss Curve')
plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()
```




Accuracy: 97.08%

Classification Report:

	precision	recall	f1-score	support
0	0.95	1.00	0.97	159
1	1.00	0.94	0.97	149
accuracy			0.97	308
macro avg	0.97	0.97	0.97	308
weighted avg	0.97	0.97	0.97	308





KNN USING KERAS

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from keras.models import Sequential
from keras.layers import Dense
import seaborn as sns

# Load the Heart dataset
df = pd.read_csv("heart1.csv") # Replace with actual file path or URL

# Preprocess the data (e.g., fill missing values, select features, etc.)
df.fillna(df.median(), inplace=True) # Fill missing values with median (for simplicity)

# Select features and target (assuming 'target' is the column with labels)
X = df.drop('target', axis=1).values # Features (adjust column name)
y = df['target'].values # Target

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Standardize the features (important for neural networks)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Create the ANN model using Keras
model = Sequential()

# Add input layer (with 30 units) and first hidden layer (with 64 units)
model.add(Dense(64, input_dim=X_train_scaled.shape[1], activation='relu'))

# Add second hidden layer (with 32 units)
model.add(Dense(32, activation='relu'))

# Add output layer (binary classification with 1 unit and sigmoid activation)
model.add(Dense(1, activation='sigmoid'))

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

